

```
In [2]: # Import necessary libraries
import pandas as pd
from sklearn.datasets import fetch_openml
from sklearn.preprocessing import StandardScaler
from sklearn.cluster import KMeans
import matplotlib.pyplot as plt
import seaborn as sns

# Step 1: Load the Dataset
# Fetch the Wine Quality dataset and save it as a CSV file
data = fetch_openml(name='wine-quality-red', as_frame=True)
df = pd.concat([data.data, data.target.rename('quality')], axis=1)
df.to_csv('wine_quality.csv', index=False) # Save as CSV

# Load the CSV file
df = pd.read_csv('wine_quality.csv')

# Print the dataset structure to verify column names
print(df.head())
print(df.columns)

# Step 2: Preprocess the Data
# Separate features and target variable
features = df.drop(columns=['quality'])
target = df['quality']

# Standardize the features
scaler = StandardScaler()
features_scaled = scaler.fit_transform(features)

# Step 3: Apply KMeans Clustering
# Initialize KMeans with 3 clusters
kmeans = KMeans(n_clusters=3, random_state=42, n_init=10)
kmeans.fit(features_scaled)

# Add cluster labels to the dataframe
df['cluster'] = kmeans.labels_

# Step 4: Visualize the Clusters
# Use the correct column names from the dataset
# Replace 'fixed acidity' and 'alcohol' with the actual column names
if 'fixed acidity' in features.columns and 'alcohol' in features.columns:
    x_col = 'alcohol'
    y_col = 'fixed acidity'
else:
    # Replace with alternative columns if those do not exist
    x_col = features.columns[0]
    y_col = features.columns[1]

plt.figure(figsize=(10, 6))
sns.scatterplot(
    x=features[x_col],
    y=features[y_col],
    hue=df['cluster'],
    palette='viridis',
    s=100,
    alpha=0.7
)
plt.title(f"Clustering of Wine Quality Data ({x_col} vs. {y_col})")
plt.xlabel(x_col.capitalize())
plt.ylabel(y_col.capitalize())
plt.legend(title="Cluster")
plt.grid(True)
plt.show()
```

C:\Users\Kritika Agrahari\anaconda3\Lib\site-packages\sklearn\datasets_openml.py:1002: FutureWarning: The default value of `parser` will change from `liac-arff` to `auto` in 1.4. You can set `parser='auto'` to silence this warning. Therefore, an `ImportError` will be raised from 1.4 if the dataset is dense and pandas is not installed. Note that the pandas parser may return different data types. See the Notes Section in fetch_openml's API doc for details.

```
warn(
    fixed_acidity  volatile_acidity  citric_acid  residual_sugar  chlorides  \
0                7.4                0.70         0.00             1.9         0.076
1                7.8                0.88         0.00             2.6         0.098
2                7.8                0.76         0.04             2.3         0.092
3               11.2                0.28         0.56             1.9         0.075
4                7.4                0.70         0.00             1.9         0.076

    free_sulfur_dioxide  total_sulfur_dioxide  density    pH  sulphates  \
0                   11.0                   34.0  0.9978  3.51         0.56
1                   25.0                   67.0  0.9968  3.20         0.68
2                   15.0                   54.0  0.9970  3.26         0.65
3                   17.0                   60.0  0.9980  3.16         0.58
4                   11.0                   34.0  0.9978  3.51         0.56

    alcohol  quality
0         9.4         5
1         9.8         5
2         9.8         5
3         9.8         6
4         9.4         5
Index(['fixed_acidity', 'volatile_acidity', 'citric_acid', 'residual_sugar',
      'chlorides', 'free_sulfur_dioxide', 'total_sulfur_dioxide', 'density',
      'pH', 'sulphates', 'alcohol', 'quality'],
      dtype='object')
```

